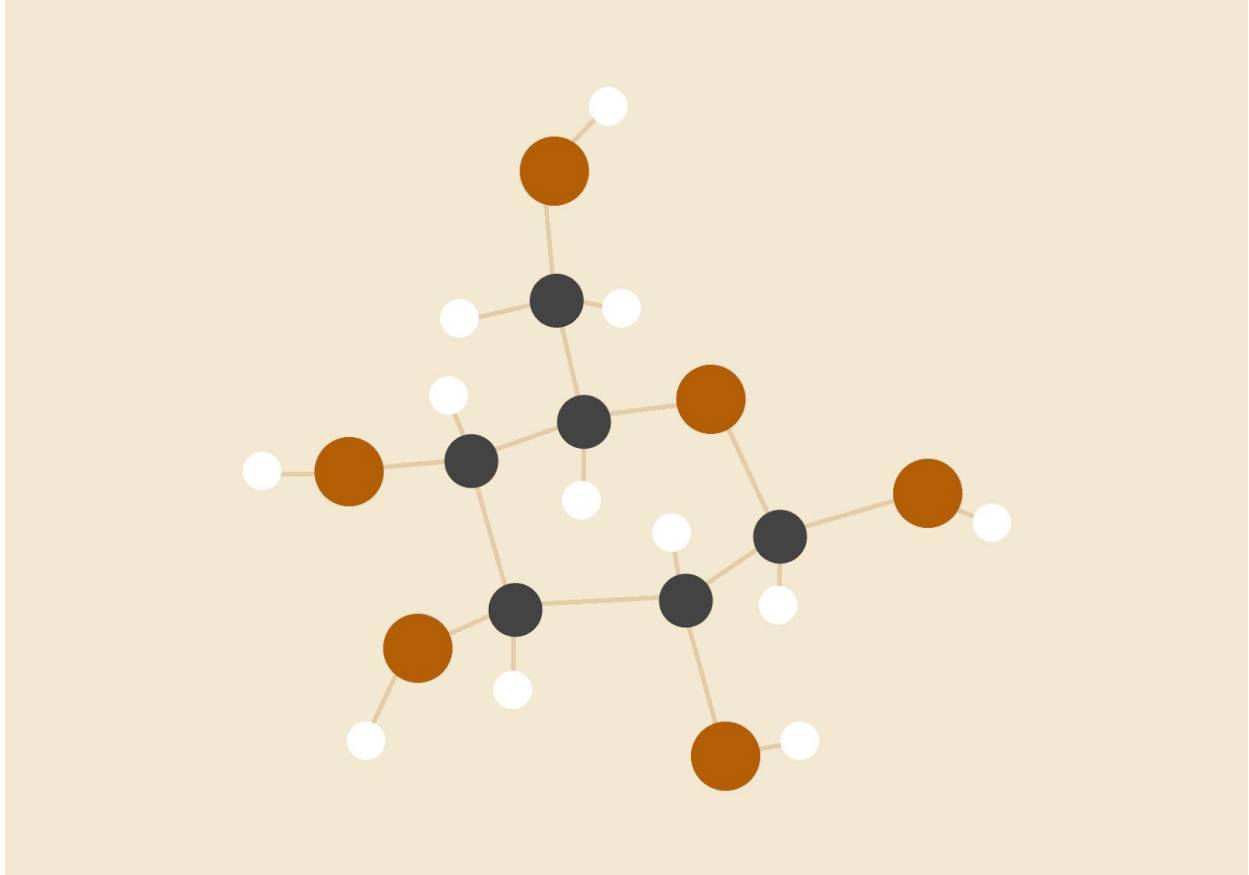


EE514 ASSIGNMENT

Data Analysis and Machine Learning



Niall Lyons

7.07.2020

Student Number: 13493628

EE514

Table Of Contents

Table Of Contents	1
DATA SET DESCRIPTION	2
RELATED WORK	3
Pre Processing and Cleaning	3
INSPECTION	3
CLEANING	4
SPLITTING THE DATA	4
FEATURE EXTRACTION	5
VERIFICATION	6
EXPLORATORY DATA ANALYSIS	6
SUPERVISED CLASSIFICATION	13
CLASSIFICATION ALGORITHMS	18
HOLDOUT VALIDATION	18
LOGISTIC REGRESSION	18
GAUSSIAN NB	19
MULTINOMIAL NB	19
K-FOLD VALIDATION	20
CLASSIFICATION MODELS K-FOLD COMPARISON TABLE	22
CLASSIFICATION MODELS HOLDOUT VALIDATION COMPARISON TABLE	22
ENSEMBLE METHODS	23
HOLDOUT VALIDATION	23
ADABOOST CLASSIFIER	23
GRADIENT BOOSTING CLASSIFIER	24
RANDOM FORESTS CLASSIFIER	25
EXTRA TREES CLASSIFIER	26
K-FOLD VALIDATION	27
ENSEMBLE MODELS HOLDOUT VALIDATION COMPARISON TABLE	28
ENSEMBLE MODELS K-FOLD COMPARISON TABLE	29
MODEL SELECTION	29
RANDOMFORESTCLASSIFER	30
LOGISTIC REGRESSION	31
GAUSSIAN NB	33
TUNED MODEL COMPARISON TABLE	33
MODEL EVALUATION	34
DISCUSSION AND CONCLUSIONS	35
REFERENCES	36

DATA SET DESCRIPTION

In the machine learning and the natural language processing space, sarcasm detection is a relevant and widely researched area. Many researchers work on this problem by scraping tweets collected using hashtags, labels and replies to tweets. Each of these datasets, although large and contain plenty of data, can be quite noisy. With tweets, comes misspellings, informal language and hate speech resulting in noisy datasets. In this report I will use a dataset specifically designed to overcome the disadvantages of Twitter datasets. It is a Headlines dataset for sarcasm detection that is collected from two news websites. These two websites, namely TheOnion and the HuffPost specialise in producing sarcastic and non sarcastic headlines respectively. The advantage of this dataset over previous datasets is that the headlines are written by professionals with no spelling mistakes using formal language. TheOnion's sole purpose is to create sarcastic headlines resulting in high quality data with less noise.

Each record in the dataset has three attributes:

- `is_sarcastic`: 1 if the record is sarcastic, otherwise 0
- `Headline`: the headline of the news article
- `Article_link`: link to the original news article

The dataset statistics can be seen in the below table. This shows us that we have a balanced dataset, with a very small margin of bias toward non sarcastic records. However, this is not a cause of concern, as it is so small.

Statistic/Dataset	Headlines
# Records	28,619
# Sarcastic records	13,635
# Non-sarcastic records	14,984

Fig 1. Dataset Statistics

RELATED WORK

The main work that has been done on this dataset is using a Hybrid Neural Network.[7] Mainly CNN and LSTM architectures have been found to improve the baseline by 5%. Given the scale and popularity of this area I feel that there is a lot more to come in this area. NLP researchers are keen to understand how sarcasm can be detected and LSTM and BLSTM architectures are finding those answers.

Pre Processing and Cleaning

Throughout this project the approach to cleaning the data was simple, which included a sequence of three steps aiming at producing high quality data to produce the best quality models. These steps included:

1. Inspection: Detect what data was given and how it could be adjusted and enhanced
2. Cleaning: Fix anomalies in the data
3. Splitting of data: Splitting data to training and test sets
4. Feature Extraction: Using TF-IDF and CountVectorizer
5. Verifying: Inspect results of data after cleaning to ensure correctness

INSPECTION

The data is in JSON format and was easily loaded by parsing the file. I was then able to pass the parsed file to a pandas dataframe to inspect what data was given. I started by looking at the head of the data, which can be seen in figure two. As explained in the description it has 3 attributes. I then wanted to find out if we had any missing data in each attribute by using `df.isna().sum()` and found that all values were present. I looked through some of the links aligned to each headline to understand what headline represented and how “sarcastic” was personified. Although interesting, for the case of this project it is not needed.

	is_sarcastic	headline	article_link
0	1	thirtysomething scientists unveil doomsday clo...	https://www.theonion.com/thirtysomething-scienc...
1	0	dem rep. totally nails why congress is falling...	https://www.huffingtonpost.com/entry/donna-edw...
2	0	eat your veggies: 9 deliciously different recipes	https://www.huffingtonpost.com/entry/eat-your-...
3	1	inclement weather prevents liar from getting t...	https://local.theonion.com/inclement-weather-p...
4	1	mother comes pretty close to using word 'strea...	https://www.theonion.com/mother-comes-pretty-c...

Fig 2. Head of data

CLEANING

The first thing to do was to delete the `article_link` column from the dataframe as it is obsolete. Although each headline is written by professionals it is not noise free. As we are dealing with words, the first method I wrote was to make all words lowercase. As with sentences comes different punctuations, which would hamper our models. Each headline was cleaned of all punctuations and anomalies. As we now have a clean sentence I wanted to clean them up some more. I researched both lemmatization and stemming and came to the conclusion that lemmatization in this case is better as it performs morphological analysis of the words. Lemmatization is the process of grouping together the different inflected forms of a word so they can be analysed as a single item, bringing a context to the words. [1] I now have data that can be understood, however I still have useless words such as 'the', 'is', 'a' that give no context. These are called stop words and have little input to our models on taking up processing time. These words are removed from all headlines. The data is now cleaned and can be split up into train and test sets.

SPLITTING THE DATA

Before splitting the data into train and test sets, I shuffled the data to ensure that there is absolutely no bias. The data is split 75% training and 25% testing. From reading different NLP papers, it seems that this ratio split is becoming the norm with a good distribution seen in the figure below. From here the training data was stored and used for data exploration while the test data was put away to be used in later stages.

```
Train Data Distribution
0    0.524087
1    0.475913
Name: is_sarcastic, dtype: float64
Test Data Distribution
0    0.522152
1    0.477848
```

Fig 3. Data Distribution

FEATURE EXTRACTION

This element of the project is the area which will determine how good your model really is. There are many different ways to approach text related problems. For example, HashingVectorizer, Bag of Words, TFidfVectorizer, TFidfTransformer and CountVectorizer. For the purpose of this project I focused on two, TFidfVectorizer and CountVectorizer. These two are the state of the art in NLP. I started with CountVectorizer as it is the simpler of the two, here I got an idea of the data I was dealing with as a starting point. This method returns a sparse vector with the number of times a word appears in a document and uses this as its weight. However, this is quite a simple method and can run into problems resulting in skewed models. A much better approach is TFidfVectorizer which tokenises documents, learns the vocabulary and inverse document weights. TFidf stands for “Term Frequency – Inverse Document” will define how important a word is in a document, while also looking at how relevant it is to the overall document. A score for each word is assigned, the mathematics behind the TFidf scores can be found in [2]. To ensure I was happy with how TFidfVectorizer was running I tuned the hyperparameters on a trial and error basis to ensure the data would produce the best models. This trial and error stage happened when validating the model and determining its accuracy. These set up for TFidfVectorizer and CountVectorizer included,

- Encoding set to utf-8, as this is the default feature used to decode
- Decode_error is set to ‘ignore’, meaning that if a byte sequence is given to analyze that contains characters not of the given encoding it is ignored.
- Stop_words was not set as these have already been removed
- Strip_accents set to ascii, as this is the fastest mapping, and accents are removed from the data
- Analyser set to ‘word’, meaning that the features analysed are of word form
- Max_df set to 0.98, this will now build a vocabulary, not using terms that have a document frequency higher than the threshold. Trial and error was used in order to find this threshold
- Min_df set to 0.15, again this was found using trial and error. This will build a vocabulary while ignoring terms with a document frequency of lower than the threshold set.

VERIFICATION

To ensure data was to a high standard and ready to be trained, I inspected the data

looking at each data frame and was happy with how it looked. I made one method that brought together the whole pre processing stage as shown:

```
def get_data():
    input_file = './Sarcasm_Headlines_Dataset.json'
    data = list(parseJson(input_file))

    df = pd.DataFrame(data)
    df = df[['is_sarcastic', 'headline']]
    df = shuffle(df)

    df['headline'] = df.headline.apply(to_lower)
    df['headline'] = df.headline.apply(punctuation_stp_words_lemmatize)
    tfidf_vectorizer = TfidfVectorizer(analyzer='word', encoding='utf-8',
    decode_error='ignore', strip_accents='ascii')
    X = tfidf_vectorizer.fit_transform(df['headline'])
    X = pd.DataFrame(X.toarray(),
    columns=tfidf_vectorizer.get_feature_names())
    y = df['is_sarcastic']
    X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.25, random_state=42)

    return X_train, X_test, y_train, y_test
```

EXPLORATORY DATA ANALYSIS

In this section I came familiar with the data that I had and was able to visualise in my mind my end goal. Exploratory data analysis gives you an opportunity to become accustomed to your data and find what is useful and non useful. I began by analysing how many sarcastic and non sarcastic headlines I had. This showed me I had a balanced dataset which I was happy with.




Fig. 4 Total Headlines Count

From here I thought about what differences could there be between both. I looked at the number of words in each headline as well as the number of characters, to get a feeler for any trends in the classifications. I used a boxplot to gain a better understanding. A box plot is a standardised way of showing the distribution of data based across min, first quartile, median, third quartile and maximum. It's a good indication of any outliers in your data and if it is symmetrical. In figure 5, the minimum which is calculated by $Q1 - 1.5 * IQR$ and maximum calculated by $Q3 + 1.5 * IQR$ are slightly higher for the sarcastic headlines in each case. Again we can see the median of sarcastic is slightly higher also in this case. Interestingly, it can also be seen that there are more outliers in the sarcastic headlines with one especially standing out above the rest. In figure 6, it can be seen slightly the same trends however with even more outliers in the sarcastic headlines. The sarcastic headlines also seem to have a slightly larger interquartile range, however not large enough to say the data is skewed.




Fig. 5 Words in headlines




Fig. 6 Characters in headlines

From observing this information I wanted to then look at the average word lengths in each class. It can be seen in figure 7, that the majority of the words in both classes are of length four to six characters with sarcastic headlines having a slightly larger kernel density.




Fig. 7 Average word length in headlines

From here I wanted to see how badly stop words affected the data and in figure 8, it can be clearly seen that in both classes the frequency of stop words is high. Meaning that keeping them would severely affect the performance of the models. Code for this part of the project is in the notebook, ML_classification_pipeline.




Fig. 8 Non-sarcastic (Left), Sarcastic (Right)

I then removed the stop words and applied TfidfVectorizer and plotted the top 20 most frequent words. I split the data into sarcastic, non sarcastic and both classes using data

frames to examine how different the word frequencies were in each. Figure 9, shows all of the headlines, where figure 10 and 11 show sarcastic and non-sarcastic words respectively.




Fig.9 Charts of Top 20 Words of all Headlines (TFidf)

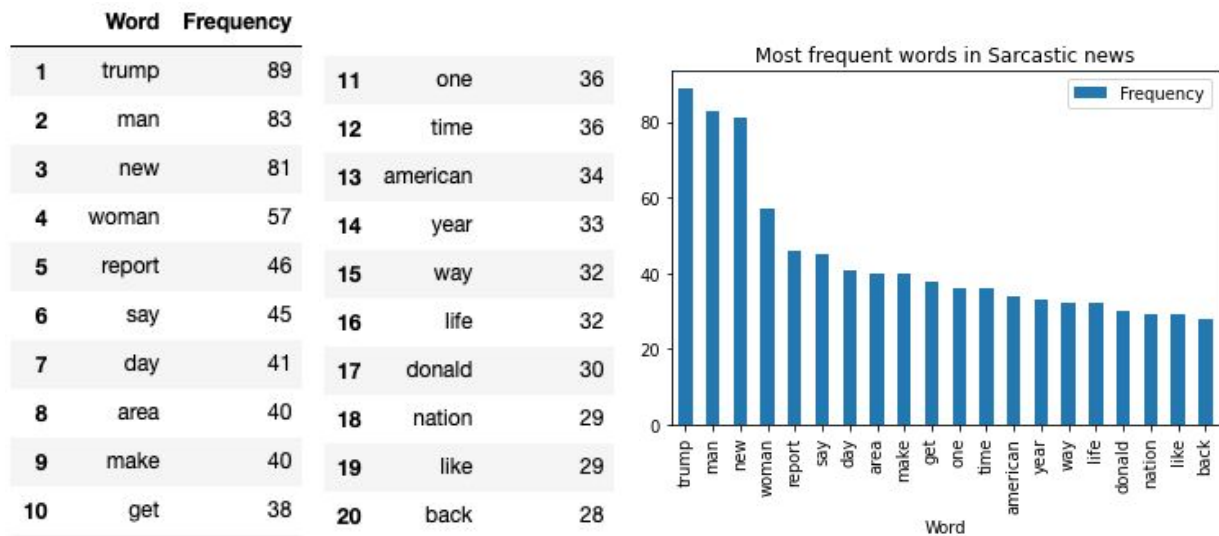


Fig. 10 Charts of Top 20 Words Sarcastic Headlines (TFidf)

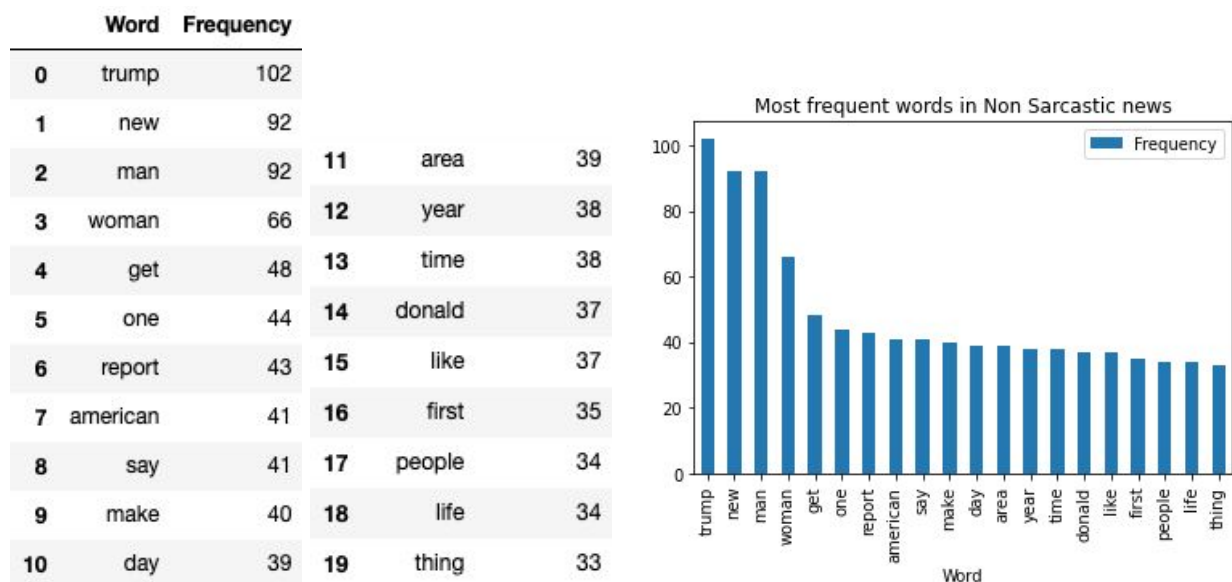


Fig. 11 Charts of Top 20 Words Non Sarcastic Headlines (TFidf)

To compare and contrast both TfidfVectorizer and CountVectorizer I re-applied the method using CountVectorizer. Figure 12, shows all of the headlines, where figure 13 and 14 show sarcastic and non-sarcastic words respectively. Code for this part of the project is in the notebook, CountVectorizer_Analysis.

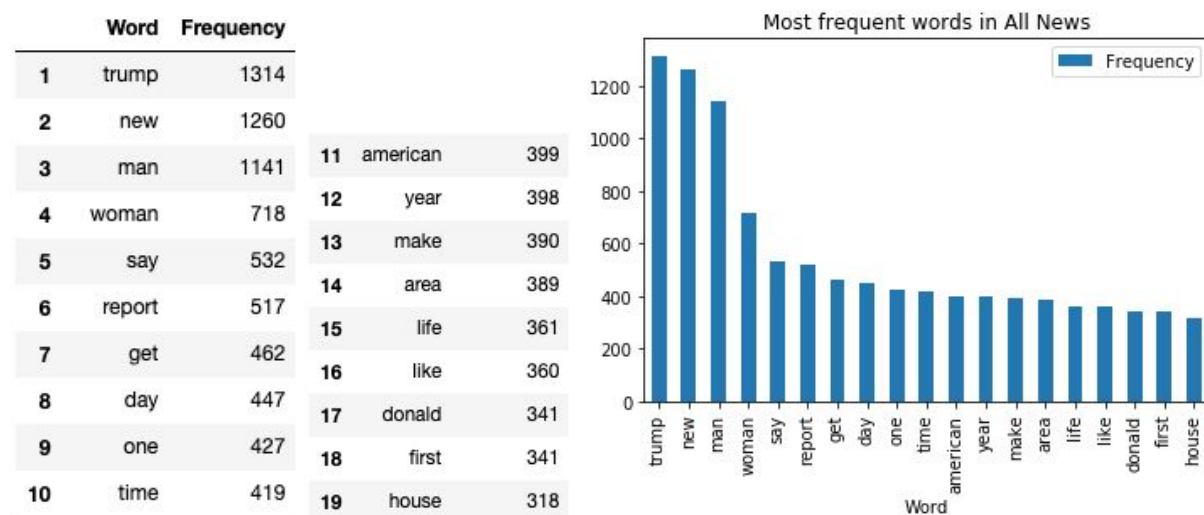


Fig. 12 Charts of Top 20 Words Headlines (CountVectorizer)

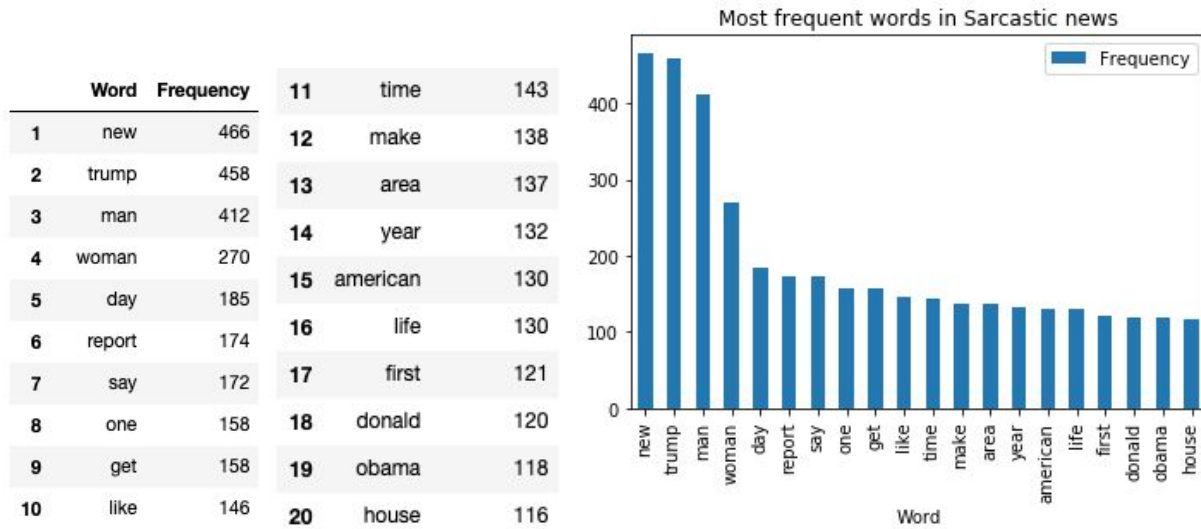


Fig. 13 Charts of Top 20 Words Sarcastic Headlines (CountVectorizer)




Fig. 14 Charts of Top 20 Words Non Sarcastic Headlines (CountVectorizer)

From generating these plots it can be seen that the same vocabulary is used when using TF-idf and CountVectorizer, however in some cases the weights of certain words vary. This proves the point that TFidf is a better approach and from this I was then able to move forward with the TFidf dataset.

In conclusion, I gained valuable insights into my data and learned alot about both classes and how they were shaped. TF-idf will be used for training, however I will also

experiment with CountVectorizer on the test set. The exploration stage of a machine learning process is vital to ensure that the right data is passed to the model for training.

SUPERVISED CLASSIFICATION

In this section I look at the different types of supervised classification methods that can be used to produce different models. The training and accuracies of the models are looked at in detail in this section. I looked at a number different models using both a holdout validation approach and the k-fold cross validation approach. I picked an array of linear and non linear classification algorithms namely (Classification_Models Notebook):

1. Logistic Regression
2. K-Nearest Neighbour
3. Decision Tree Classifier
4. Gaussian Naive Bayes
5. Multinomial Naive Baye
6. Support Vector Classifier

I then went on to investigate ensemble boosting and bagging methods (Ensemble_Models Notebook). Namely,

1. AdaBoost Classifier
2. GradientBoosting Classifier
3. Random Forest Classifier
4. ExtraTrees Classifier

For all classification models I compared both holdout validation and K-Fold validation methods. This means that in the holdout validation method I used the test set on the trained model to validate the performance of the model. The data is split 75:25 training and test respectively. This technique can be useful, however it can lead to model overfitting and underfitting. However, this problem is rectified using K-fold cross validation. Here, the data is divided into k-folds and trained on k-1 folds with one fold held for testing. This method is usually preferred as you can train on multiple trains and test splits, which in turn gives you a better outlook on how the model will perform on unseen data. However, this is computationally more expensive to run and can take a lot of power and time to run on large datasets.

Throughout the course of this analysis I will be using confusion matrices, classification

reports, precision recall curves, AUC scores and ROC curves to analyse each model.

A confusion matrix is shown in figure 15. Each row in a confusion matrix represents an actual class, while each column represents a predicted class.

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN	TN

Fig.15 Confusion Matrix

This matrix is extremely useful for accessing a classification model. True Positives, False Positives, True negatives and False Negatives. These represent:

- True Positives (TP): Predicted positive and it's true. In our case the model predicted a sarcasm headline is sarcastic and that is true.
- False Positives (FP): Predicted positive and it's false. In our case the model predicted a sarcasm headline is sarcastic but it is actually not.
- True negatives (TN): Predicted negative and it's true. In our case the model predicted a sarcasm headline is not sarcastic and it is actually not.
- False Negatives (FN): Predicted negative and it's false. In our case the model predicted a sarcasm headline is not sarcastic but it actually is.

From the confusion matrix, we move onto a classification report, which measures the quality of predictions made by the model. Looking at how many are true and how many are false. More specifically, True Positives, False Positives, True negatives and False Negatives are used to compute the metrics. Figure 16, illustrates the difference between precision and recall as shown in the classification report.




Fig. 16 Precision and accuracy

From this image it is shown that it's imperative that a model has high accuracy and high precision. Precision is the accuracy of positive predictions the classifier makes. It is defined as:

$$\text{precision} = \frac{TP}{TP + FP}$$

Where TP is the number of true positives and FP is the number of false positives. Precision is widely used beside recall this is the ratio of positive instances that are correctly detected by the classifier. It is defined as:

$$\text{recall} = \frac{TP}{TP + FN}$$

Where FN is the number of false negatives. The accuracy of a model is also given in the classification report, this is a performance measure that is a ratio between corrected

predicted cases and overall total cases. Having a high accuracy does not mean your model is very good, it depends on the shape of your data and also if FP and FN outcomes are close together. Accuracy is defined as:

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fp + fn}$$

The F1 score of a model is also outlined in the report also. This is a combination of both precision and recall and is useful to compare classifiers. It is the harmonic mean of precision and recall, meaning that it gives more weight to lower values. It is defined as:

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{TP}{TP + \frac{FN+FP}{2}}$$

To understand these metrics even deeper I also developed a precision recall curve. This shows the trade off between precision and recall for different thresholds. The higher the area under the curve the higher the recall and precision where high precision relates to a low false positive rate, and high recall relates to a low false negative rate. The higher the score the higher the accuracy. Figure 17, shows a perfect precision recall curve. However, if you were to get this after training your model there is a good chance something is wrong!




Fig. 17 Ideal Precision Recall Curve

Another metric I used to calculate the models is a receiver operating characteristic (ROC) curve. Here the true positive rate (TPR) is plotted against the false positive rate (FPR), which is the amount of negative observations that are incorrectly classified. In the ROC curve there is a trade off, the higher the TPR the more false positives the classifier produces. From this curve, we can also compare classifiers using the metric area under curve (AUC). The best AUC score will be 1.0.

In real world machine learning projects, to determine whether to use a precision recall (PR) curve or an ROC curve. Most of the time ROC curves are used, however if we want to specifically look at false positives instead of false negatives the PR curve would be used. Figure 18 shows different ROC curves and their corresponding AUC score.




Fig. 18 ROC Curves and corresponding AUC scores

CLASSIFICATION ALGORITHMS


```
LR_model.fit(X_train, y_train)
GNB_model.fit(X_train, y_train)
NB_model.fit(X_train, y_train)
SVM_model.fit(X_train, y_train)
```

HOLDOUT VALIDATION


LOGISTIC REGRESSION

	precision	recall	f1-score	support
0	0.79	0.83	0.81	3757
1	0.80	0.76	0.78	3398
accuracy			0.80	7155
macro avg	0.80	0.79	0.79	7155
weighted avg	0.80	0.80	0.80	7155

```
[[3119 638]
 [ 824 2574]]
LR accuracy score: 0.7956673654786862
LR precision score: 0.8013698630136986
LR recall score: 0.757504414361389
LR average precision-recall score: 0.72220542959365
```



AUC: 0.8739055352512078




GAUSSIAN NB


	precision	recall	f1-score	support
0	0.63	0.83	0.72	3794
1	0.71	0.45	0.55	3361
accuracy			0.66	7155
macro avg	0.67	0.64	0.64	7155
weighted avg	0.67	0.66	0.64	7155

[[3166 628]
[1840 1521]]

GNB accuracy score: 0.6550663871418588
GNB precision score: 0.7077710563052583
GNB recall score: 0.4525438857482892
GNB average precision-recall score: 0.5774602872411



AUC: 0.6436858209700811

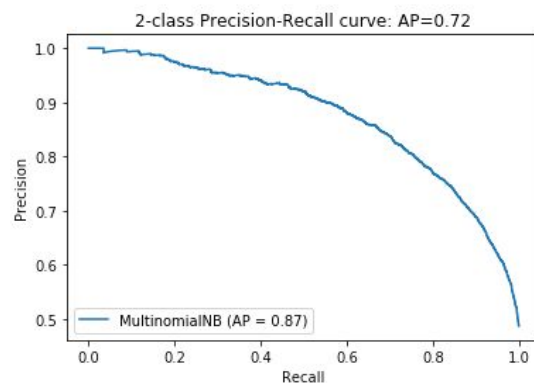


MULTINOMIAL NB


	precision	recall	f1-score	support
0	0.78	0.85	0.82	3794
1	0.81	0.74	0.77	3361
accuracy			0.80	7155
macro avg	0.80	0.79	0.79	7155
weighted avg	0.80	0.80	0.80	7155

[[3226 568]
[887 2474]]

NB accuracy score: 0.7966457023060797
NB precision score: 0.8132807363576594
NB recall score: 0.7360904492710503
NB average precision-recall score: 0.7226174348801

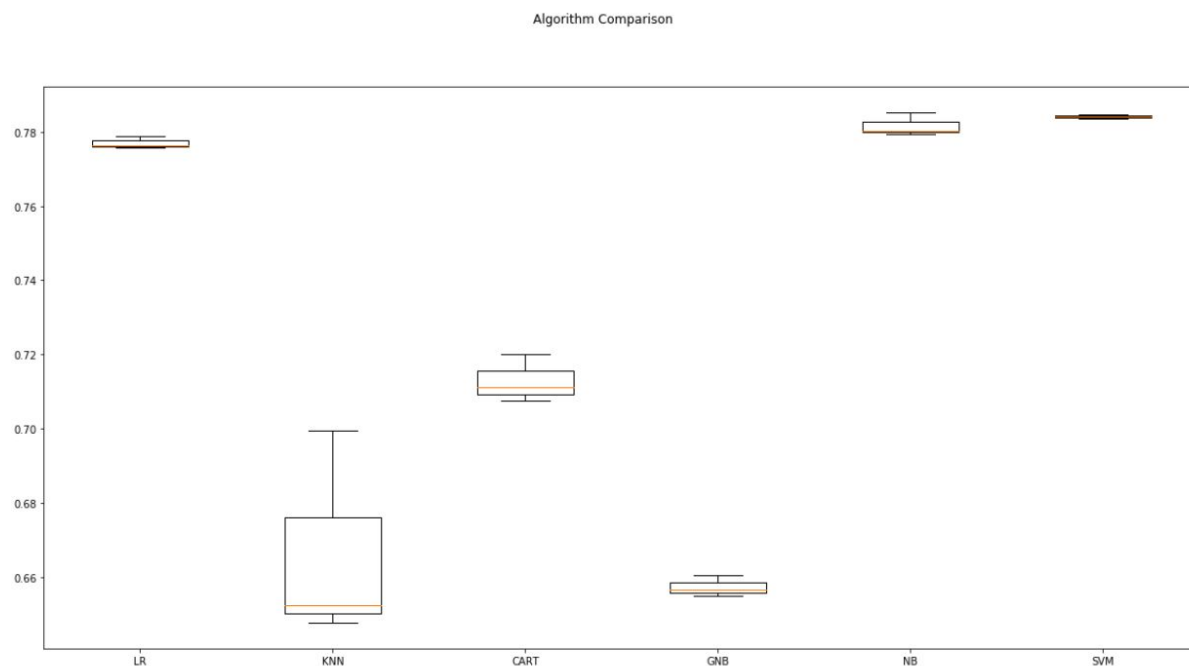


AUC: 0.8812989770565876



K-FOLD VALIDATION

In K-Fold validation, looking at the distribution of accuracy values calculated across all folds is vital to ensure you understand the outcome of your model. The below plot shows tight distributions across Logistic regression , Naives Baye and SVC which is good as it suggests low variance. However, Gaussian Naive Baye being performing so poor is slightly surprising.



As the attributes are not varied very much, apart from KNN there is no need to standardize the data in any way. By training the models and using `cross_val_score` with `k` is 3 the resulting mean of all results and std of results were:

```
LR
LR: 0.776975 (0.001367)
KNN
KNN: 0.666650 (0.023320)
CART
CART: 0.712915 (0.005253)
GNB
GNB: 0.657426 (0.002296)
NB
NB: 0.781588 (0.002561)
SVM
SVM: 0.783964 (0.000502)
```

SVC model performs the best over 3 folds with a very low STD. However, very closely behind are logistic regression and Naive Bayes. Both are computationally more efficient. The snippet of code to produce this outcome:

```
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
```

```

models.append(('SVM', SVC()))
results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=num_folds, random_state=seed)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold,
    scoring=scoring) results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

```

CLASSIFICATION MODELS K-FOLD COMPARISON TABLE

K-Fold Validation (K=3)		
Classifier	Accuracy(%)	STD
Logistic Regression	78%	0.001367
KNN	67%	0.023320
Decision Tree Classifier	71%	0.005253
Gaussian NB	66%	0.002296
Multinomial NB	78%	0.002561
Support Vector Classifier	78%	0.000502

CLASSIFICATION MODELS HOLDOUT VALIDATION COMPARISON TABLE

Holdout Validation		
Classifier	Accuracy(%)	Precision
Logistic Regression	80	0.81
KNN	67	0.82
Gaussian NB	66	0.71
Multinomial NB	81	0.81

ENSEMBLE METHODS

The main cause of error in machine learning models is noise, bias and variance. Ensemble methods try to minimise these factors. Their traits include improving the accuracy and overall performance of the models. As I am not overly familiar with these methods I wanted to try them out.

The below classifiers and analyses specifically use holdout validation to analyse the performance of the model.

HOLDOUT VALIDATION

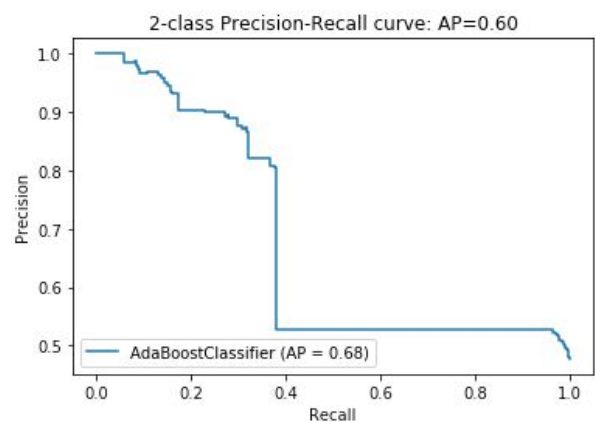
ADABOOST CLASSIFIER

An AdaBoost classifier is a meta-estimator that begins by fitting a classifier on the original dataset and then fits additional copies of the classifier on the same dataset but where the weights of incorrectly classified instances are adjusted such that subsequent classifiers focus more on difficult cases. [3]


	precision	recall	f1-score	support
0	0.62	0.92	0.74	3746
1	0.81	0.38	0.52	3409
accuracy			0.66	7155
macro avg	0.71	0.65	0.63	7155
weighted avg	0.71	0.66	0.63	7155

[[3434 312]
[2115 1294]]


```
AB accuracy score: 0.6607966457023061
AB precision score: 0.8057285180572852
AB recall score: 0.3795834555588149
AB average precision-recall score: 0.60143869940311
```



AUC: 0.7074883591485557



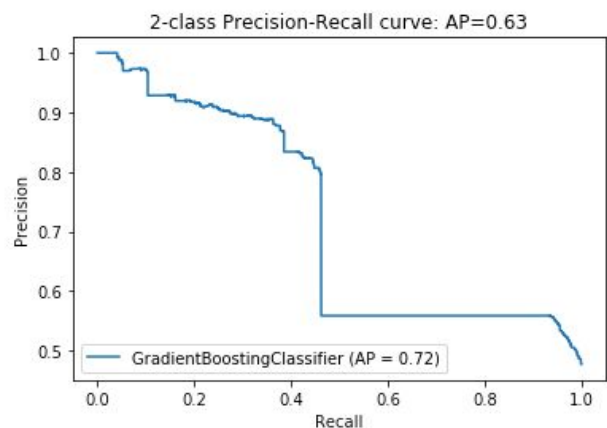
GRADIENT BOOSTING CLASSIFIER

GB builds an additive model in a forward stage-wise fashion; it allows for the optimization of arbitrary differentiable loss functions. In each stage n_{classes} regression trees are fit on the negative gradient of the binomial or multinomial deviance loss function. Binary classification is a special case where only a single regression tree is induced. [4]


	precision	recall	f1-score	support
0	0.63	0.95	0.75	3746
1	0.87	0.38	0.53	3409
accuracy			0.68	7155
macro avg	0.75	0.66	0.64	7155
weighted avg	0.74	0.68	0.65	7155

```
[[3553 193]
 [2118 1291]]
```

```
GBM accuracy score: 0.6770090845562544
GBM precision score: 0.8699460916442049
GBM recall score: 0.37870343209152246
GBM average precision-recall score: 0.6254683421287
```



AUC: 0.7502858236034542




RANDOM FORESTS CLASSIFIER

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. The sub-sample size is controlled with the `max_samples` parameter if `bootstrap=True` (default), otherwise the whole dataset is used to build each tree. [5]


	precision	recall	f1-score	support
0	0.75	0.85	0.80	3746
1	0.81	0.68	0.74	3409
accuracy			0.77	7155
macro avg	0.78	0.77	0.77	7155
weighted avg	0.78	0.77	0.77	7155

```
[[3201 545]
 [1086 2323]]
```

```
RF accuracy score: 0.7720475192173305
RF precision score: 0.8099721059972106
RF recall score: 0.681431504840129
RF average precision-recall score: 0.7037224817181
```



AUC: 0.854005375363133



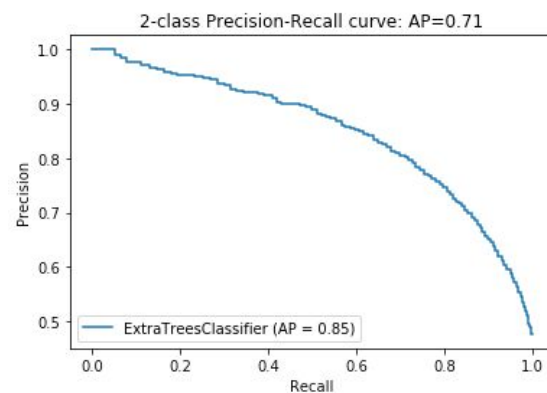
EXTRA TREES CLASSIFIER

This class implements a meta estimator that fits a number of randomized decision trees (a.k.a. extra-trees) on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. [6]


	precision	recall	f1-score	support
0	0.75	0.86	0.80	3746
1	0.82	0.68	0.74	3409
accuracy			0.78	7155
macro avg	0.78	0.77	0.77	7155
weighted avg	0.78	0.78	0.77	7155

[[3240 506]
[1096 2313]]

ET accuracy score: 0.7761006289308177
 ET precision score: 0.8205037247250798
 ET recall score: 0.6784980932824876
 ET average precision-recall score: 0.70988980744617




AUC: 0.8577697896823787



K-FOLD VALIDATION

In this box plot I compare the accuracies of the different ensemble methods. Just like above as the attributes are not varied very much, there is no need to standardize the data in any way.



By training the models and using `cross_val_score` with `k = 3` the resulting mean of all results and std of results were:

```

AB
AB: 0.645546 (0.003833)
GBM
GBM: 0.665812 (0.002889)
RF
RF: 0.756663 (0.004622)
ET
ET: 0.763744 (0.003514)

```

The code was to produce these results:

```

ensembles = []
ensembles.append(('AB', AdaBoostClassifier()))
ensembles.append(('GBM', GradientBoostingClassifier()))
ensembles.append(('RF', RandomForestClassifier()))
ensembles.append(('ET', ExtraTreesClassifier()))

results = []
names = []

for name, model in ensembles:
    kfold = KFold(n_splits=num_folds, random_state=seed)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold,
    scoring=scoring)
    results.append(cv_results)
    names.append(name) msg = "%s: %

```

ENSEMBLE MODELS HOLDOUT VALIDATION COMPARISON TABLE

Holdout Validation				
Classifier	Accuracy	Precision	AUC	Misclassified Headlines
AdaBoost Classifier	66%	0.80	0.70	9731
Gradient Boost Classifier	68%	0.87	0.75	9159
Random Forest Classifier	77%	0.81	0.86	6583
ET Classifier	78%	0.82	0.85	6297

ENSEMBLE MODELS K-FOLD COMPARISON TABLE

K-Fold Cross validation (K=3)		
Classifier	Accuracy (%)	STD
AdaBoost Classifier	65%	0.003833
Gradient Boost Classifier	67%	0.002889
Random Forest Classifier	76%	0.004622
Extra Trees Classifier	76%	0.003514

Interesting from the above results the Extra trees classifier has performed the best over both validation techniques with a precision of 0.82 and AUC of 0.85. When starting the ensemble methods I had thought that Adaboost would have been the most significant algorithm for this problem but I was proven wrong.

MODEL SELECTION

At this point I have engineered a lot of models, using different validation techniques. Different data has been passed and I am now happy that I decided on TfidfVectorizer as it yielded the best results. As of now the model hyperparameters were chosen through trial and error and defaults. Hyperparameters are set by the developer and can have a large impact on the outcome of a model and will in most cases make the model perform to a higher standard. In order to get the best hyperparameters for a model we can use GridSearchCV which searches for the best hyperparameters from values that are passed to fine tune the model and raise the robustness and performance. For example, `n_estimators` in `RandomForestClassifier` are often tuned using `GridSearchCV`. However, this method of fine tuning is computationally expensive and takes time to run. I found and used `RandomSearchCV`, which is computationally less expensive and yields the same results in most cases. From the code snippet below it was very important to use the hyperparameters of `RandomizedSearchCV` to my advantage. For example setting `n_jobs` to `-1`, utilizes all processors and runs jobs in parallel, without it tuned models would have taken a long time to run.

```
grid = RandomizedSearchCV(model, param_distributions=param_grid,
scoring=scoring, cv=2, n_jobs=-1, verbose=10)
```

I have chosen three different models to fine tune and look into.


RANDOMFORESTCLASSIFIER

The tuned model using RandomSearchCV yielded the following results:


```
Best: 0.760529 using {'n_estimators': 100, 'max_features': 'log2'}
0.720183 (0.004193) with: {'n_estimators': 10, 'max_features': 'auto'}
0.736908 (0.008899) with: {'n_estimators': 30, 'max_features': 'auto'}
0.741101 (0.010296) with: {'n_estimators': 50, 'max_features': 'auto'}
0.745527 (0.005404) with: {'n_estimators': 100, 'max_features': 'auto'}
0.710492 (0.003261) with: {'n_estimators': 10, 'max_features': 'log2'}
0.745527 (0.000559) with: {'n_estimators': 30, 'max_features': 'log2'}
0.757780 (0.004333) with: {'n_estimators': 50, 'max_features': 'log2'}
0.760529 (0.003634) with: {'n_estimators': 100, 'max_features': 'log2'}
```

	precision	recall	f1-score	support
0	0.75	0.90	0.82	3746
1	0.86	0.67	0.75	3409
accuracy			0.79	7155
macro avg	0.80	0.78	0.78	7155
weighted avg	0.80	0.79	0.79	7155
[[3362 384]				
[[1123 2286]]				

```
Best_RF accuracy score: 0.7893780573025856
Best_RF precision score: 0.8561797752808988
Best_RF recall score: 0.6705778820768554
Best_RF average precision-recall score: 0.73108839!
```



AUC: 0.8811124160677031



LOGISTIC REGRESSION

Across both k-fold and holdout validation Logistic regression has been performing well. I wanted to try it with many hyperparameters so I set it to:


```
penalty = ['l1', 'l2']
c = np.logspace(-4, 4, 20)
solver = ['liblinear', 'lbfgs', 'newton-cg']
param_grid = dict(C=c, penalty=penalty, solver=solver)

model = LogisticRegression()
grid = RandomizedSearchCV(model, param_distributions=param_grid,
scoring=scoring, cv=2, n_jobs=-1, verbose=10)
grid_result = grid.fit(X_train, y_train)
```


This yielded the below results. I really wanted to look into the different solvers used as I thought liblinear would fit well for this and I wanted to gain a deeper understanding of l1 and l2 penalties.

```
Best: 0.770360 using {'solver': 'liblinear', 'penalty': 'l2', 'C': 4.281332398719396}
0.768729 (0.001584) with: {'solver': 'newton-cg', 'penalty': 'l2', 'C': 11.288378916846883}
0.705507 (0.002562) with: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 0.08858667904100823}
nan (nan) with: {'solver': 'lbfgs', 'penalty': 'l1', 'C': 0.615848211066026}
0.770360 (0.000140) with: {'solver': 'liblinear', 'penalty': 'l2', 'C': 4.281332398719396}
nan (nan) with: {'solver': 'newton-cg', 'penalty': 'l1', 'C': 545.5594781168514}
nan (nan) with: {'solver': 'newton-cg', 'penalty': 'l1', 'C': 29.763514416313132}
0.739890 (0.000140) with: {'solver': 'liblinear', 'penalty': 'l1', 'C': 78.47599703514607}
0.734998 (0.000839) with: {'solver': 'liblinear', 'penalty': 'l1', 'C': 10000.0}
nan (nan) with: {'solver': 'lbfgs', 'penalty': 'l1', 'C': 78.47599703514607}
nan (nan) with: {'solver': 'lbfgs', 'penalty': 'l1', 'C': 545.5594781168514}
```

When evaluated the tuned logistic regression model



AUC: 0.8745138390044676



GAUSSIAN NB

A classifier not entirely used in industry, but I thought it would be good to research it and look into it. The code was set up as below,

```
nb_classifier = GaussianNB()
var_smoothing = np.logspace(0, -9, num=100)
param_grid = dict(var_smoothing=var_smoothing)
gs_NB = RandomizedSearchCV(nb_classifier,
                           param_distributions=param_grid,
                           cv=2,
                           n_jobs=-1,
                           verbose=10,
                           scoring='accuracy')

gs_NB.fit(X_train, y_train);
```

This yielded the below results. By changing the var_smoothing, this is the portion of the largest variance of all features that is added to variances for calculation stability.

```

Best: 0.747484 using {'var_smoothing': 0.1873817422860384}
0.660362 (0.002516) with: {'var_smoothing': 1.873817422860383e-06}
0.661619 (0.002749) with: {'var_smoothing': 4.328761281083053e-06}
0.658964 (0.002423) with: {'var_smoothing': 8.111308307896872e-07}
0.747484 (0.000093) with: {'var_smoothing': 0.23101297000831597}
0.747484 (0.000559) with: {'var_smoothing': 0.1873817422860384}
0.660781 (0.002935) with: {'var_smoothing': 2.310129700083158e-06}
0.657240 (0.003401) with: {'var_smoothing': 2.848035868435805e-09}
0.657193 (0.003448) with: {'var_smoothing': 2.310129700083158e-09}
0.657147 (0.003494) with: {'var_smoothing': 6.579332246575682e-09}
0.673779 (0.001584) with: {'var_smoothing': 0.0004328761281083057}

```

TUNED MODEL COMPARISON TABLE

Tuned Models			
Classifier	Accuracy	Precision	AUC
Random Forest Classifier	79%	0.86	0.88
Logistic Regression	79%	0.80	0.87
Gaussian NB	75%	0.85	0.86

It can be seen that the random forest classifier has a 2% increase in accuracy, a 0.05 increase in precision and 0.02 increase in AUC score. This means that by fine tuning the model is now performing better. With more time and more parameters passed, and more computational power all hyperparameters could be accessed and a much more accurate and robust model can be developed.

One model I really wanted to fine tune is the SVC, however I don't have the computational power at the moment and it is taking a very long time to do a thorough evaluation and analysis unfortunately.

MODEL EVALUATION

The final model was saved with pickle and loaded to test against the holdout validation set. The testing dataset was set up the same way as all other TfidfVectorizer sets. The below are the statistics for the final model.



	precision	recall	f1-score	support
0	0.75	0.90	0.82	3746
1	0.86	0.67	0.75	3409
accuracy			0.79	7155
macro avg	0.80	0.78	0.78	7155
weighted avg	0.80	0.79	0.79	7155


```
[[3362 384]
 [1123 2286]]
```



```
Best_RF accuracy score: 0.7893780573025856
Best_RF precision score: 0.8561797752808988
Best_RF recall score: 0.6705778820768554
Best_RF average precision-recall score: 0.731088391
```

AUC: 0.8811124160677031



The out of sample error here is 0.211 which was achieved by calculating 1 - accuracy. This metric helps you understand the fit of the model. As the out of sample error gets closer to 1, our main priority is not to overfit the data. Thus, with the out of sample error increasing and degrees of freedom increasing as previously stated overfitting can occur.

The ROC curve of the final model is good for a college project, but at a production level the AUC score would need to be much closer to 1, resulting in a lot more accurate predictions.

Overall I am happy with the final model. However, I am aware that the Naives Baye model had the highest accuracy of 80%. There are no real hyperparameters to tune in Naives bayes and that is the reason I left it as it was and explored different classifiers to enhance my learning.

DISCUSSION AND CONCLUSIONS

Doing this project encouraged me deeper to understand the fundamentals of machine learning and the process of developing an end to end project. If this was a project for production I would have continued evaluating and fine tuning each model I created to ensure that performance and robustness were of high quality. The logistics regression and SVC would have been of significant importance as I feel with some shifting of hyperparameters a lot could be done with it.

Something that I would like to do in the future is to use deep learning on this problem to enhance the performance of models while also trying new things with this very interesting dataset. In particular using Long Short Term Memory and Bidirectional LSTM networks would really enhance the performance of these models on the dataset and with some good engineering bring accuracies up to the low nineties at least.

I learned a lot doing this project, and can definitely bring the knowledge I gained into some of my job interviews. I think what I learned the most is that developing a end to end machine learning project is a long but enjoyable task with many roadblocks that you must adapt and overcome.

I found this project, and EE514 a very knowledgeable experience and I look forward to more machine learning projects in the future.

REFERENCES

- [1] NLTK, <http://www.nltk.org/api/nltk.stem.html?highlight=lemmatizer>
- [2] Toward Data Science, <https://towardsdatascience.com/tf-idf-explained-and-python-sklearn-implementation-b020c5e83275>
- [3] Sklearn, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>
- [4] Sklearn, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>
- [5] Sklearn, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>
- [6] Sklearn, <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>
- [7] Cornell University, <https://arxiv.org/abs/1908.07414>, 2019